AI evaluation | Evidence and limitations

Juniper Vale fictional work sample | October 2026 | Public verification note

What the case study contains

An independent fictional use-case design for service exception note triage. The proposed 150-scenario test mix is 100 standard, 20 ambiguous, 15 adversarial and 15 safety/high-impact examples. A published case-study PDF reports illustrative baseline and controlled metric values.

What could not be verified

The original project package contains a 19-column AI evaluation CSV header followed by five incompatible three-field staff-training rows. It is not a usable 150-case dataset. Therefore the published aggregate scenario scores cannot be independently recalculated from the case-level inputs and outputs.

What has been provided instead

AI_Evaluation_Scenario_Summary.csv accurately transcribes the five published synthetic scenario metrics and labels them non-reproducible. AI_Evaluation_Test_Plan.csv documents the intended test groups and controls. Neither file is represented as case-level model execution evidence.

What would complete this demonstration

Create and retain 150 synthetic notes and expected labels, execute the defined baseline and controlled prompts against a specified model/version, save outputs and reviews, then calculate metrics with a reproducible script. Until that happens, the model scores should be read as scenario assumptions, not verified AI benchmark results.