AI Evaluation | Evidence Boundary

Juniper Vale independent fictional case | October 2026 | Public evidence note

Available evidence

The original project contains a proposed test mix of 100 standard, 20 ambiguous, 15 adversarial and 15 safety/high-impact scenarios. The planned controls include human review, instruction-boundary checks and fallback.

Unverified historical scores

Earlier versions of this case study displayed aggregate modeled percentages. The original evaluation CSV has an incompatible header and data rows, so those percentages cannot be reproduced from retained individual test examples. They have been removed from the featured page and updated case study.

What is provided now

The downloadable AI_Evaluation_Test_Plan.csv describes categories, intended checks and expected controls. It is not an executed test log. No model execution, completed pilot, or historical benchmark is asserted.

What verification would require

Save test notes, expected labels, precise model/version and prompt configuration, the complete actual outputs, adjudication decisions and an executable metrics calculation. Without those, avoid numeric effectiveness claims.

Jeff Pursley | https://jeffpursley-coder.github.io/ | October 2026