Responsible AI Service-Note Triage

Juniper Vale fictional scenario | Design work sample, not a deployed AI system or measured benchmark

Business question

Can an AI tool help staff turn a variable service-exception note into structured suggestions without assuming authority over safety, scheduling or customer communication?

Proposed workflow

1. Read the untrusted note as data, not instructions. 2. Apply deterministic safety checks. 3. Retrieve only approved, current policy and taxonomy references. 4. Produce a constrained summary, category, escalation suggestion and follow-up type. 5. Validate the output against business rules. 6. Require authorized human review.

Where AI should not decide

Eligibility, capacity, consent, route assignment, message sending, exception closure and supervisor overrides remain rule-based or human-controlled. The design gives AI no write/action permissions.

Risk and failure controls

Untrusted text may attempt prompt injection; an allowed output schema and no action tools limit its effect. Ambiguity should trigger abstention. Missing or stale knowledge should block confident guidance. Explicit safety language triggers human review regardless of model output.

Evaluation design

The original scenario proposed 150 test cases split across standard, ambiguous, adversarial and safety-sensitive notes. Retained case-level inputs and model outputs could not be verified, so this case study claims no accuracy, recall or prompt-injection success rate. A real test would preserve the model/version, input, expected label, actual output, reviewer decision and result for every case.

Pilot and fallback decision

Before a pilot, verify the test evidence, human review controls, monitoring, escalation triggers and manual fallback. If the AI path fails, staff can continue using the existing structured taxonomy and knowledge base. No real client pilot or production result is asserted.

Jeff Pursley | https://jeffpursley-coder.github.io/ | October 2026